Avaliação de estilos de código para árvore de decisão em GPU com Microbenchmarks

Jeronimo Penha, Alysson Kelvim, Olavo Silva, Icaro Moreira, José A. Nacif, **Ricardo Ferreira**- Universidade Federal de Viçosa

CNPq
FAPEMIG
CAPES

Summary

- Random Forest and GPU
- Three Implementation: IF, **without IF**, Memory
- Which is the best ? Depends on Depth and Trees
- Conclusions

Random Forest

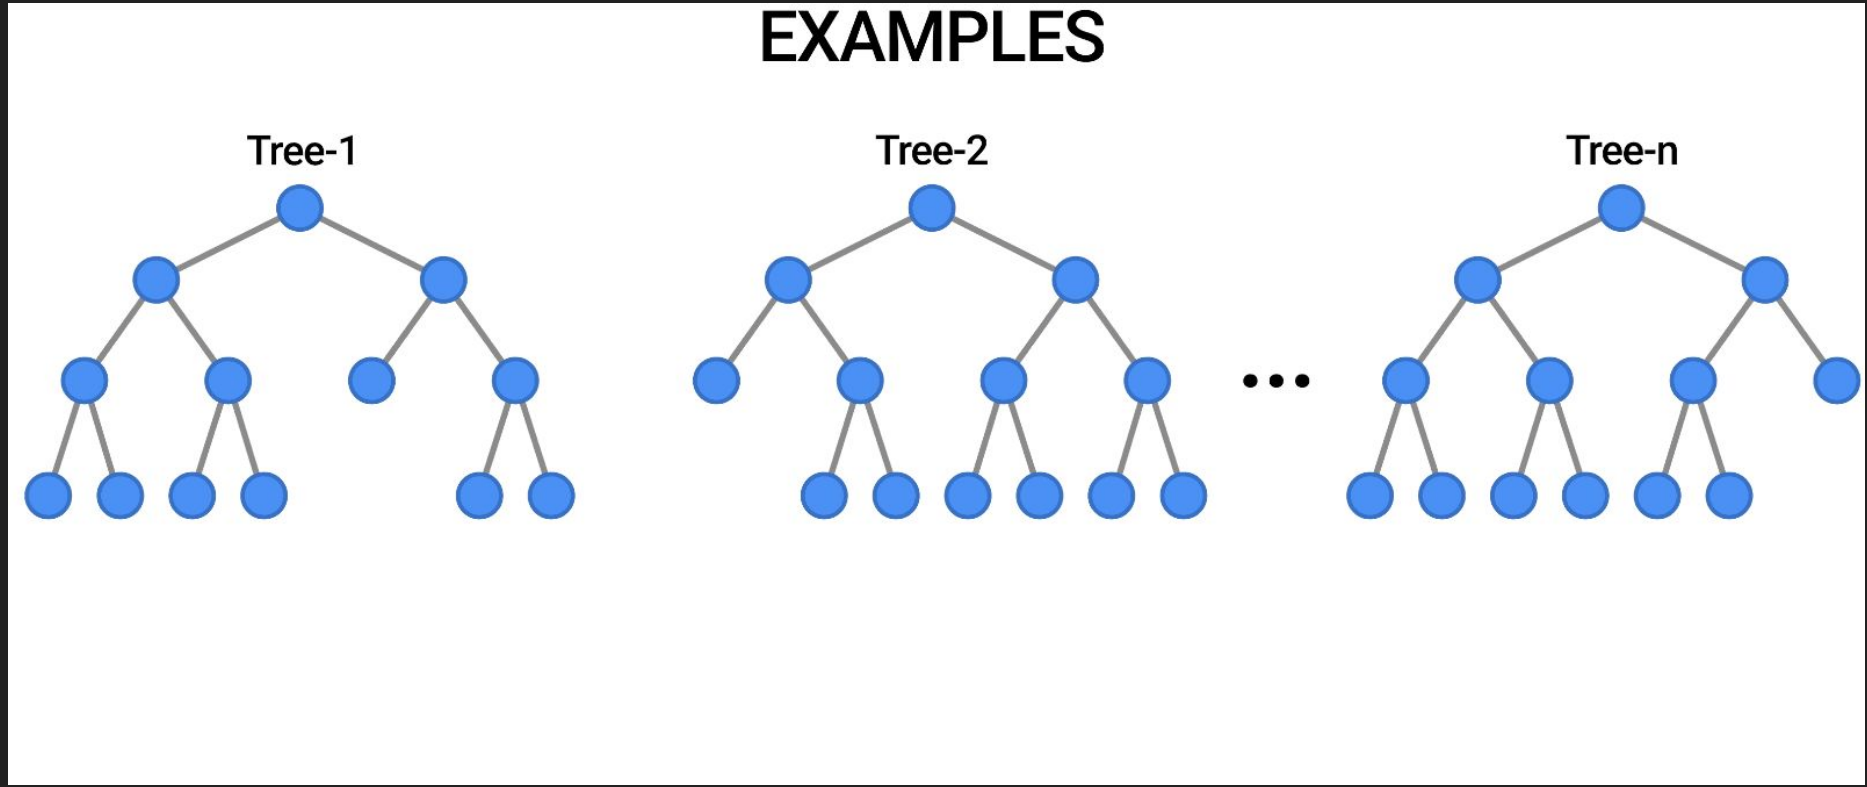

Figure from Tensor Flow Blog

Random Forest

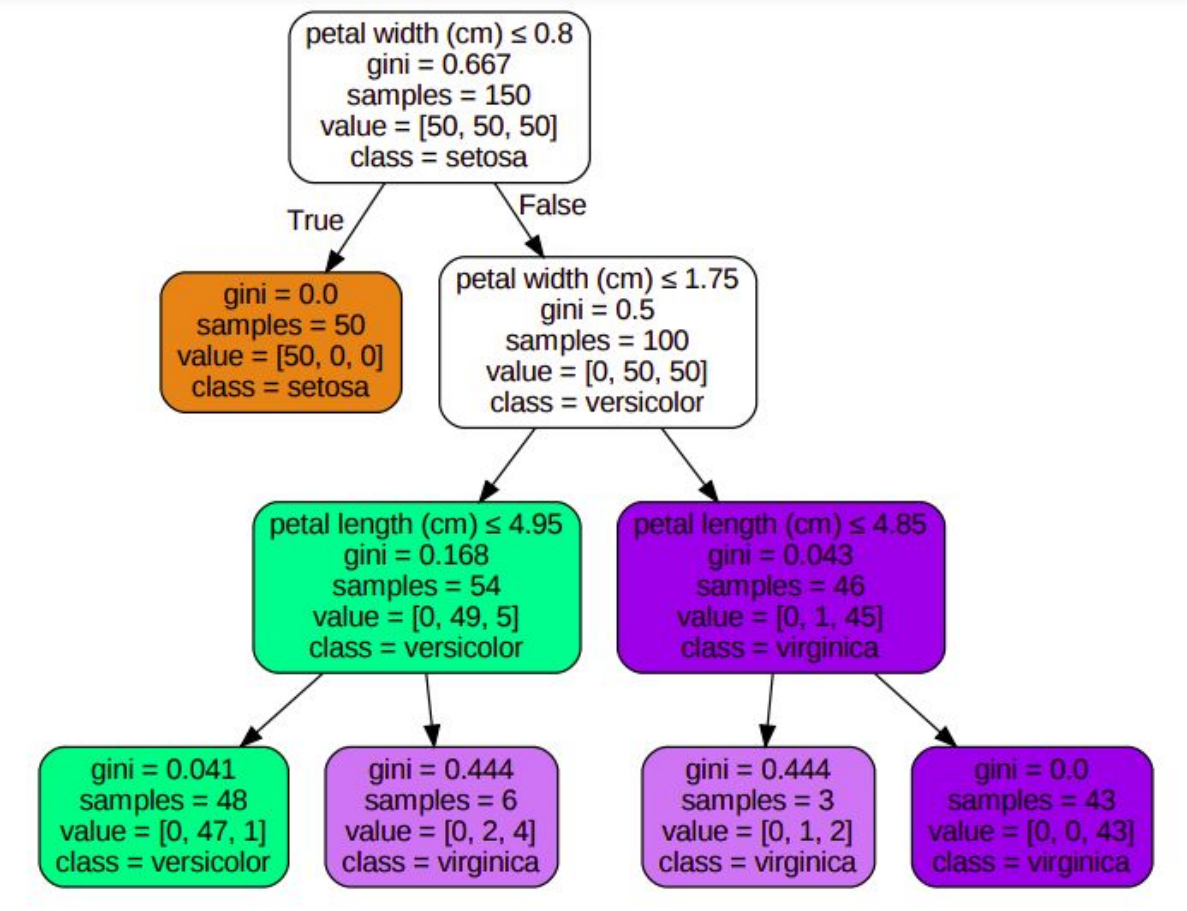


```
 if (x[3] <= 0.80)
      return 0;
else {
      if (x[3] <= 1.75) {
            if (x[2] <= 4.95)
                  return 1;
            else
                  return 2;
      } else {
            if (x[2] <= 4.850)
                  return 2;
            else
                  return 2;
      }
```

GPU and Branches


```
if (threadIdx.x < 4) {
    A;
    B;
} else {
    X;
    Y;
}
Z;
```

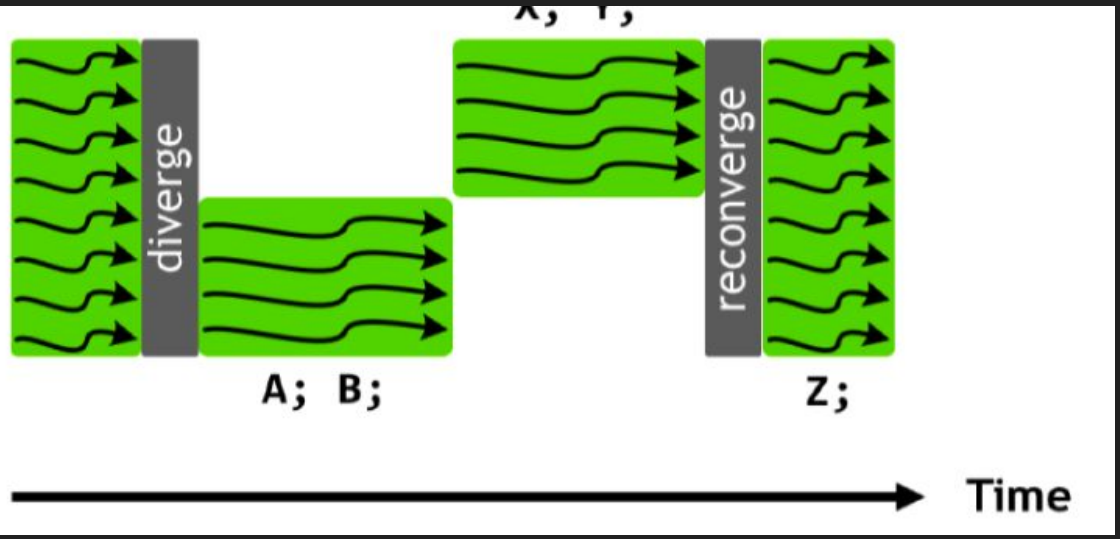

Figure extracted from here

GPU and Branches

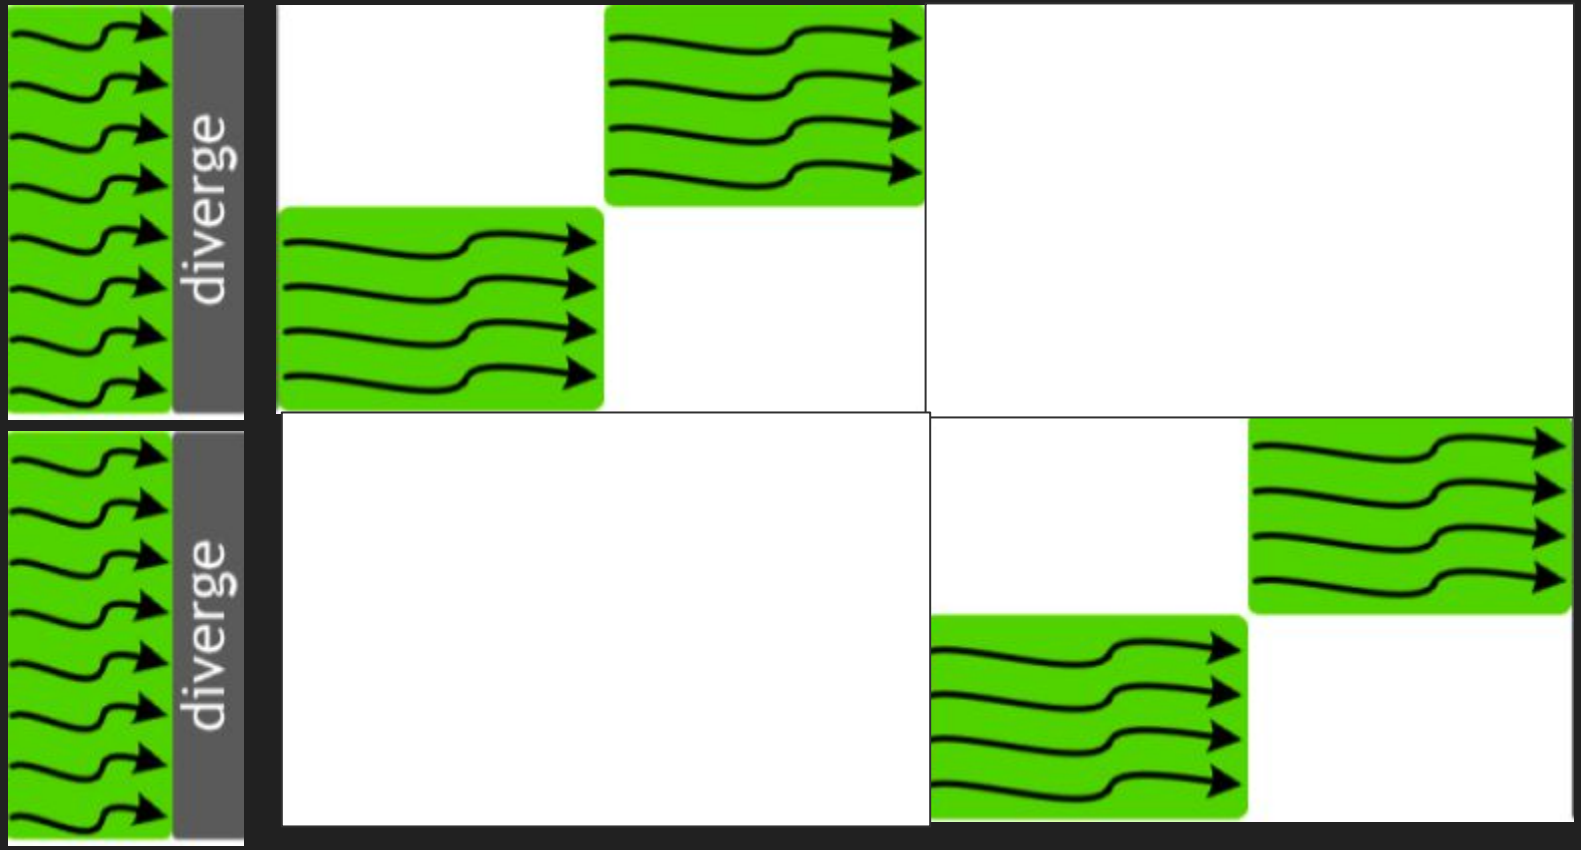

Figure adapted from here

PTX IF

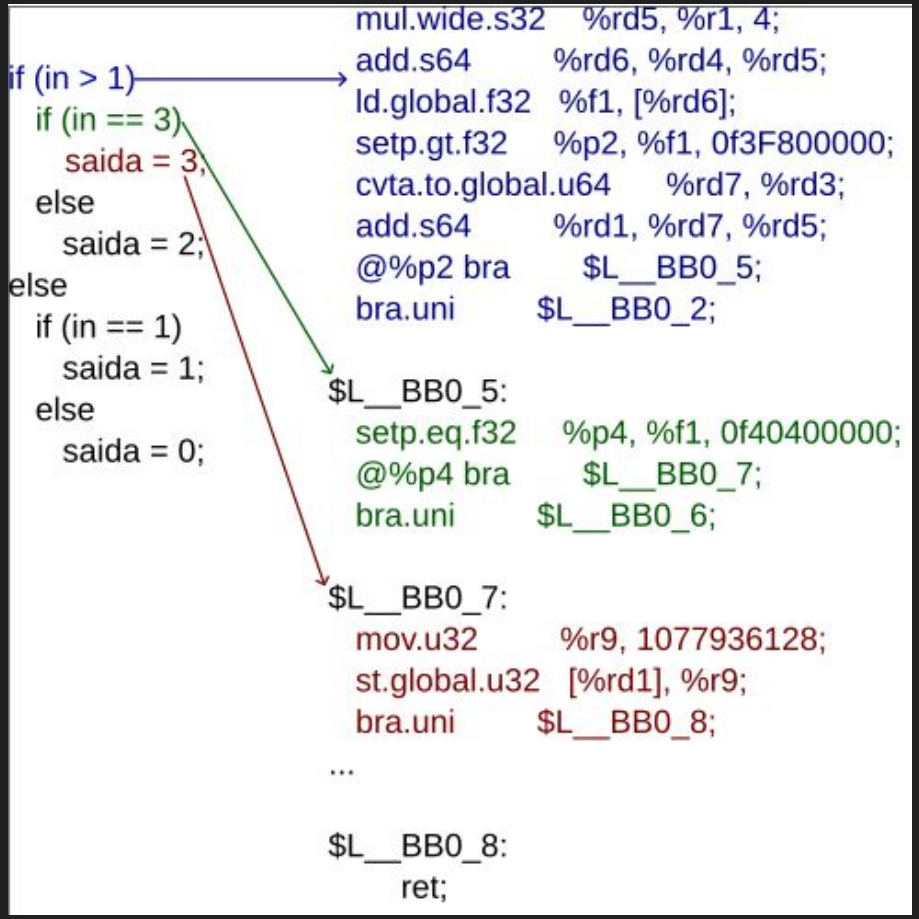

Faster approach for more than 7 levels

- no indirection
- simple branch

Research Question

- Is there an alternative approach to implementing trees without branching?

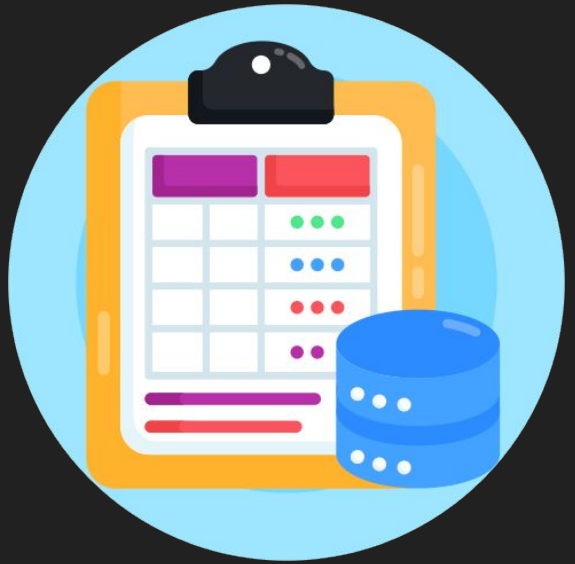

When is it advantageous to utilize a GPU?

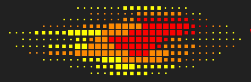

- Over 5 million registers
- Numerous memory types and units
- Thousands of computing units

Decision Tree as a table

- Used in FPGAs, No divergence, same code for all threads

	Feature	Threshold	Left	Right
0 →	Age	> 40	1	2
1	Weight	> 70	3	4
2	B. Pres	> 80	7	8

```
Index = 0
while not leaf do
        F= table.feature(index)
        data = Input(F)
        T = table.Threshold(index)
        Index = ( data > T) ?
                table.left(index):
                table.right(index);
```

Decision Tree as a table


```
Index = 0
while not leaf do
    F= table.feature(index)
    data = Input(F)
    T = table.Threshold(index)
    Index = ( data > T) ?
            table.left(index):
            table.right(index);
```

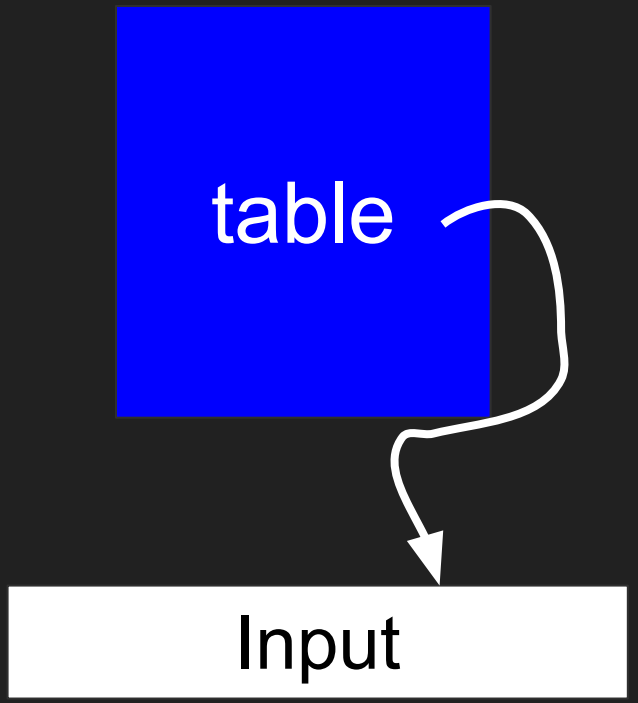

Decision Tree as a table


```
__global__ void RF(
            ...
  const float* TH,
  const int* tabela)
   {        ...        }
```

Global

```
__global__ void RF(
            ...
  float* __restrict__ TH,
  int* __restrict__ tabela)
   {        ...        }
```

read-only

```
__constant__ int tabela[TAM_TABELA];
__constant__ float TH[TAM_TH];
...
__global__ void RF(...)
   {        ...        }
```

Constant

```
__global__ void RF(
            ...
  const float* p_th,
  const int* p_tabela){

__shared__ float TH[TAM_TH];
__shared__ int tabela[TAM_TABELA];
            ...
}
```

Shared

PTX Table


```
ld.global.f32   %f3, [%rd5];
setp.lt.f32     %p2, %f3, %f2;
selp.b32        %r7, 2, 10, %p2;
ld.global.u32   %r8, [%rd6];
shr.u32         %r9, %r8, %r7;
and.b32         %r10, %r9, 255;
```

(a)

```
ld.global.nc.f32        %f3, [%rd5];
setp.lt.f32     %p2, %f3, %f2;
selp.b32        %r7, 2, 10, %p2;
ld.global.nc.u32        %r8, [%rd6];
shr.u32         %r9, %r8, %r7;
and.b32         %r10, %r9, 255;
```

(b)

```
ld.const.f32    %f3, [TH];
setp.lt.f32     %p2, %f3, %f2;
selp.b32        %r7, 2, 10, %p2;
mov.u64         %rd7, table;
ld.const.u32    %r8, [table];
shr.u32         %r9, %r8, %r7;
and.b32         %r10, %r9, 255;
```

(c)

```
ld.shared.f32   %f4, [RF::TH];
setp.lt.f32     %p4, %f4, %f3;
selp.b32        %r15, 2, 10, %p4;
mov.u32         %r16, RF)::table;
ld.shared.u32   %r17, [RF::table];
shr.u32         %r18, %r17, %r15;
and.b32         %r19, %r18, 255;
```

(d)

High latency 30 cycles

- indirection
- no divergence
- Slow approach (good for FPGA)

Three level - Global

Table in MEMORY

Trees	IF	Global	
		dir	ind
1	2,56	8,43	8,91
2	5,21	17,06	17,66
3	7,67	25,57	22,22
4	10,12	30,93	28,54

Time in Milliseconds - 50 million samples

Nvidia GTX 1070

Three level - Read Only

Table in MEMORY

Trees	IF	Read-only		Global	
		dir	ind	dir	ind
1	2,56	6,92	8,92	8,43	8,91
2	5,21	13,90	17,67	17,06	17,66
3	7,67	25,55	26,32	25,57	22,22
4	10,12	30,95	28,99	30,93	28,54

Time in Milliseconds - 50 million samples

Nvidia GTX 1070

Three level Shared

Trees	IF
1	2,56
2	5,21
3	7,67
4	10,12

Table in MEMORY

Shared		Read-only		Global	
dir	**ind**	**dir**	**ind**	**dir**	**ind**
2,53	3,70	6,92	8,92	8,43	8,91
3,20	4,96	13,90	17,67	17,06	17,66
3,96	6,50	25,55	26,32	25,57	22,22
4,65	8,08	30,95	28,99	30,93	28,54

Time in Milliseconds - 50 million samples

Nvidia GTX 1070

Three level Const

Table in MEMORY

Trees	IF	Const		Shared		Read-only		Global	
		dir	ind	dir	ind	dir	ind	dir	ind
1	2,56	2,01	2,82	2,53	3,70	6,92	8,92	8,43	8,91
2	5,21	2,40	3,78	3,20	4,96	13,90	17,67	17,06	17,66
3	7,67	2,66	6,14	3,96	6,50	25,55	26,32	25,57	22,22
4	10,12	4,67	7,75	4,65	8,08	30,95	28,99	30,93	28,54

Time in Milliseconds - 50 million samples

Nvidia GTX 1070

Three level NoIF

Table in MEMORY

Trees	IF	noIF	Const		Shared		Read-only		Global	
			dir	ind	dir	ind	dir	ind	dir	ind
1	2,56	1,99	2,01	2,82	2,53	3,70	6,92	8,92	8,43	8,91
2	5,21	2,13	2,40	3,78	3,20	4,96	13,90	17,67	17,06	17,66
3	7,67	2,41	2,66	6,14	3,96	6,50	25,55	26,32	25,57	22,22
4	10,12	2,89	4,67	7,75	4,65	8,08	30,95	28,99	30,93	28,54

Time in Milliseconds - 50 million samples

Nvidia GTX 1070

Proposal to implement Without IF

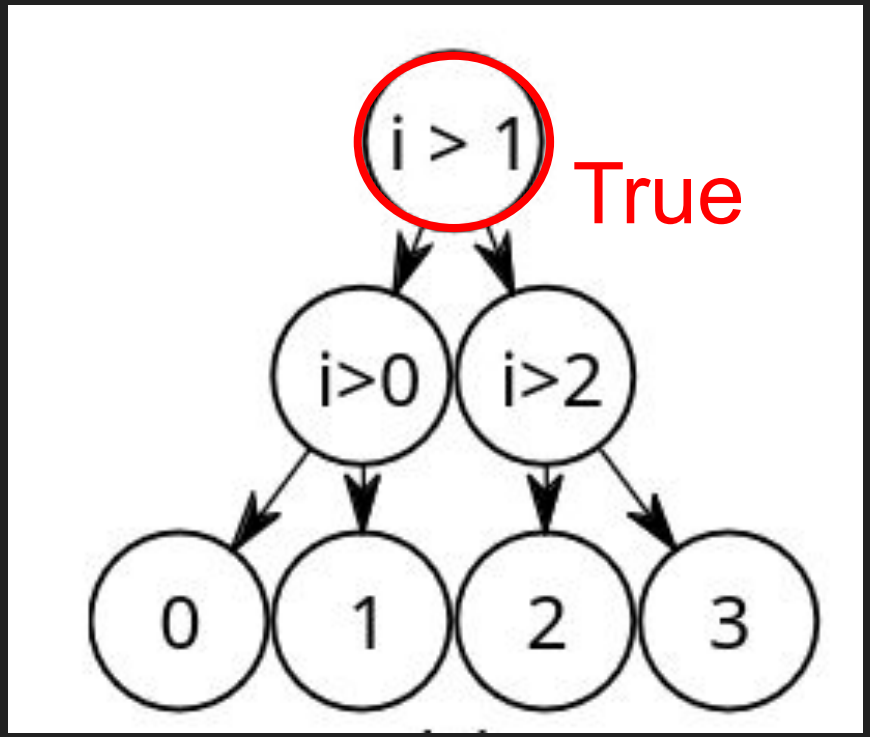

Root = (in > 1)

leaf = root * (2 + (in >2))

leaf += (1-root)*(in>0)

Output = leaf

Proposal to implement Without IF

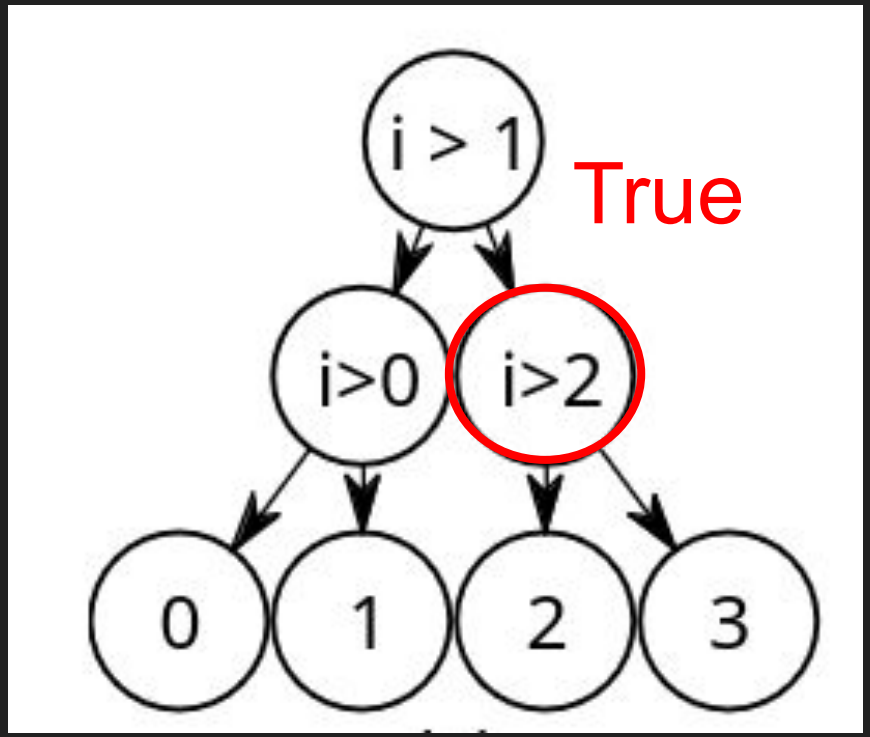

Root = (in > 1)

leaf = root * (2 + (in >2))

leaf += (1-root)*(in>0)

Output = leaf

Proposal to implement Without IF

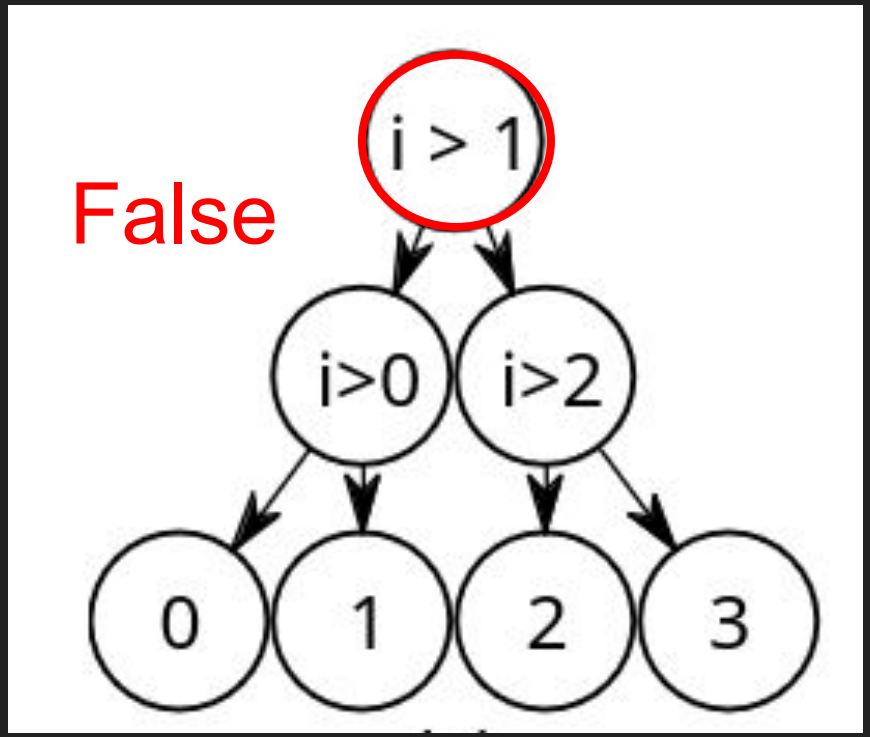

Root = (in > 1)

leaf = root * (2 + (in >2))

leaf += (1-root)*(in>0)

Output = leaf

Three Comparisons + 5 add/mult !!!

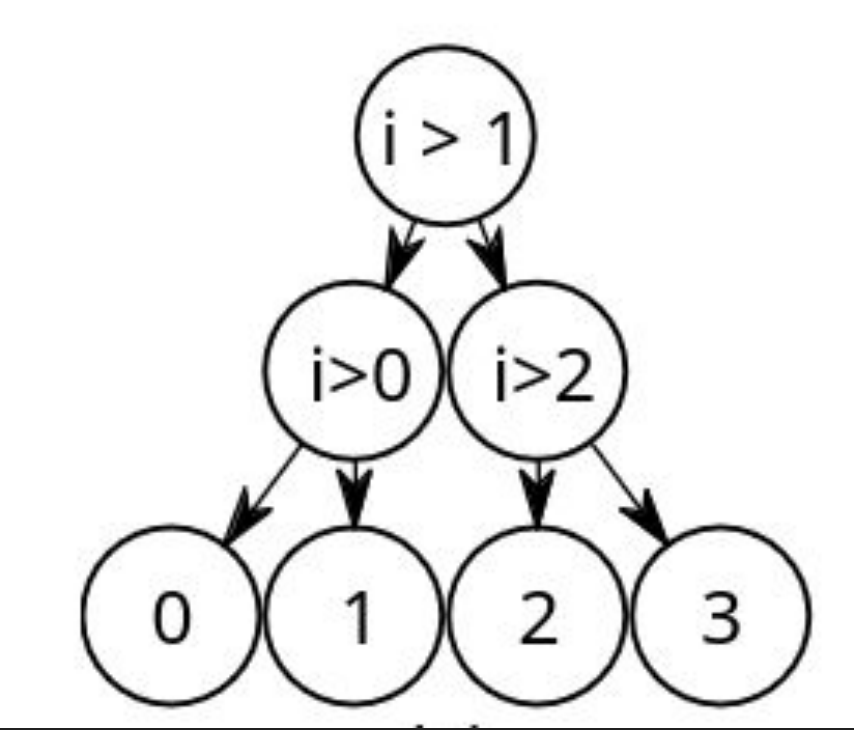

Root = (in > 1)

leaf = root * (2 + (in >2))

leaf += (1-root)*(in>0)

Output = leaf

Seven Comparisons…

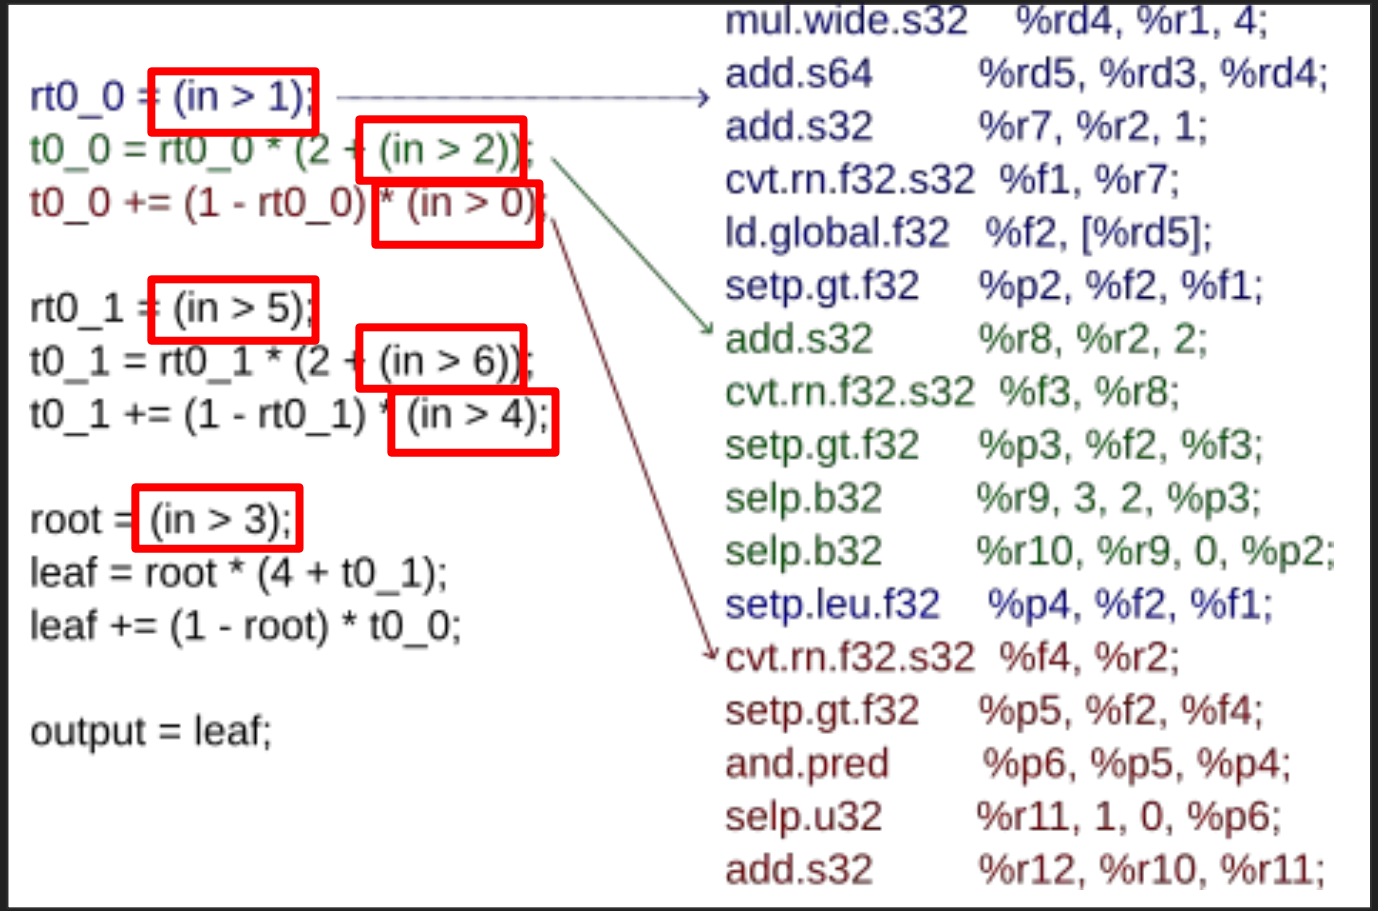

Which is better IF or No IF ?

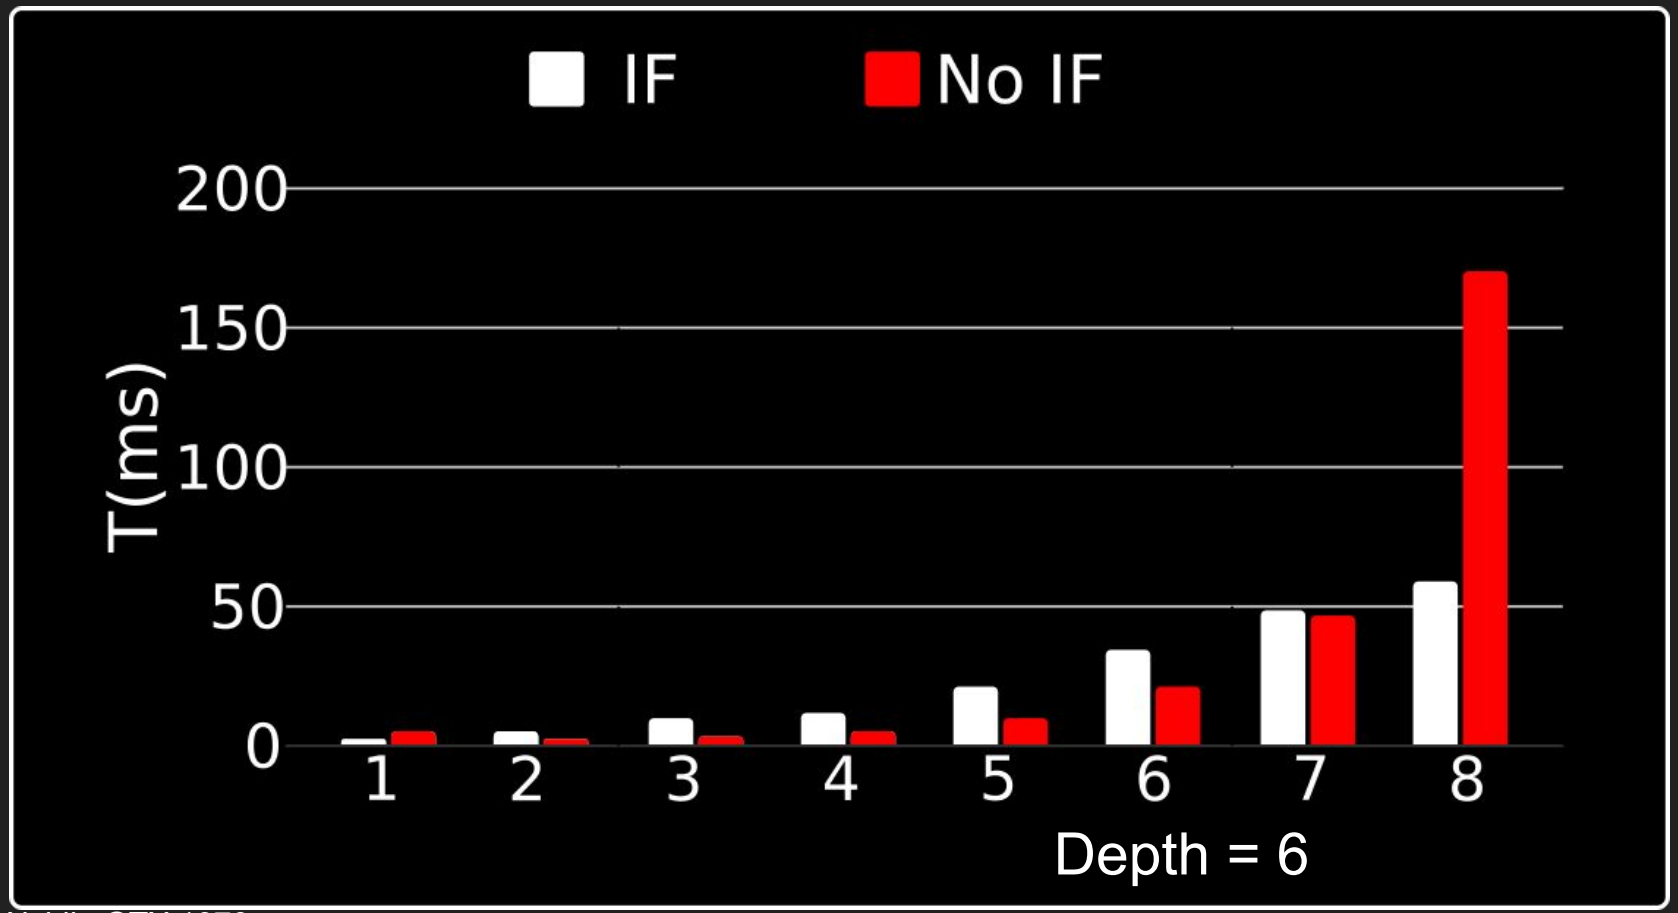

Nvidia GTX 1070

Depth == 7 ? Time in ms

Trees	IF	No IF	memory
1	10,07	13,04	13,04
2	22,07	24,18	29,46
3	35,61	34,41	43,94
4	48,74	48,34	58,43

Nvidia GTX 1070

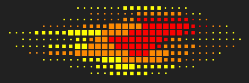

Which is better:
a few deep trees or numerous shadow trees?

IF implementation

10 Trees Depth 5 = 53.31ms 2.2x faster
5 Tress Depth 10 = 127.31

No IF
10 Trees Depth 5 = 33.4ms 3.8x faster !

Real Datasets depth=6

WSCAD 2023

No IF

IF

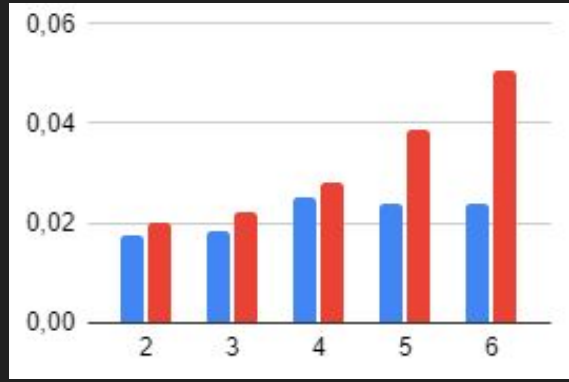

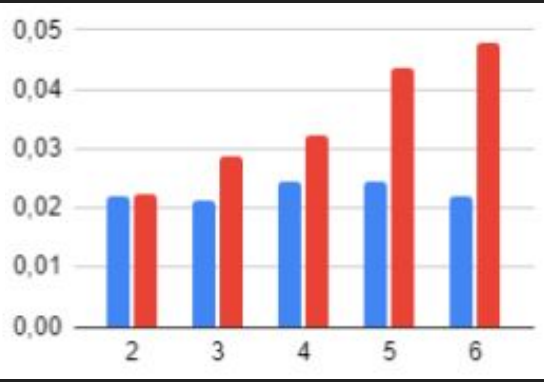

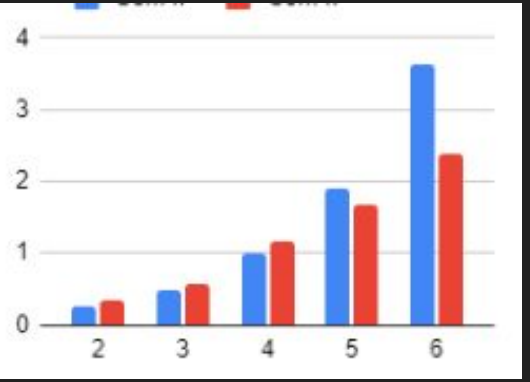

	Hospital	Adult	Susy
features	18	14	18
size	59.557	32.561	50.00.000

GPU versus OpenMP CPU

i7-7 3.6G 8 cores 16 Ths

Trees	D	NoIF		IF		Table	
	5	**3,62**	510,29	6,43	72,18	5,39	189,39
1	6	**6,56**	1257,27	9,21	75,24	7,57	221,13
	7	13,04	2598,33	**10,07**	100,57	15,14	254,60
	5	**5,61**	1235,78	10,88	109,70	8,23	362,16
2	6	**11,55**	2620,45	15,58	148,03	13,44	423,52
	7	24,18	5102,83	**22,07**	218,01	29,67	544,08
	5	**7,70**	1936,71	16,08	177,18	11,35	544,25
3	6	**15,77**	3874,22	24,30	301,82	19,69	751,58
	7	**34,41**	7601,23	35,61	326,44	44,17	865,45
	5	**10,23**	2572,75	21,32	310,98	14,32	799,97
4	6	**21,62**	5085,65	34,42	343,99	25,87	958,26
	7	**46,95**	10111,26	48,34	440,42	58,66	1078.87

GPU versus OpenMP CPU

30x i7-7 3.6G 8 cores 16 Ths

Trees	D	NoIF		IF		Table	
1	5	**3,62**	510,29	6,43	72,18	5,39	189,39
	6	**6,56**	1257,27	9,21	75,24	7,57	221,13
	7	13,04	2598,33	**10,07**	100,57	15,14	254,60
2	5	**5,61**	1235,78	10,88	109,70	8,23	362,16
	6	**11,55**	2620,45	15,58	148,03	13,44	423,52
	7	24,18	5102,83	**22,07**	218,01	29,67	544,08
GPU is		250x		15x		40x	
4	5	**10,23**	2572,75	21,32	310,98	14,32	799,97
	6	**21,62**	5085,65	34,42	343,99	25,87	958,26
	7	**46,95**	10111,26	48,34	440,42	58,66	1078.87

Conclusions

- No IF up to Depth 6
- Better many shadow trees than a few deep tree
- GPU 30x faster than CPU
- Future Work: FPGAs and compare with Real Dataset and Tree Depth+Number

Avaliação de estilos de código para árvore de decisão em GPU com Microbenchmarks

ricardo@ufv.br

Questions ?

FAPEMIG APQ-01577-22

CNPq

CAPES